

Divorce Case Recommender — Exhibition Report Data

1. Corpus size and label counts (matches the "About" tab metrics)

- **Total cases:** 56
- **Unique grounds of divorce:** 8
- **Verdict types:** 5

Grounds of divorce breakdown:

Ground	Count
Unspecified	36
Cruelty	8
Mutual Consent	3
Irretrievable Breakdown	3
Cruelty, Desertion	3
Cruelty, Irretrievable Breakdown	1
Mutual Consent, Irretrievable Breakdown	1
Desertion	1

Verdict breakdown:

Verdict	Count
Allowed	38
Dismissed	8
Disposed	6
Unknown	3
Quashed	1

Note: 36/56 cases (64%) are tagged "Unspecified" for ground of divorce — worth flagging in your write-up as a limitation of the keyword-based tagging step, and a natural place to say "future work: improve ground tagging."

2. Example query with top-5 recommendations

Query entered: *"husband repeatedly abusive, wife seeks divorce on grounds of cruelty and mental harassment"*

Case	Ground	Verdict	Similarity	Matched keywords
Raj Talreja vs Kavita Talreja (24 Apr 2017)	Cruelty	Allowed	0.180	wife, cruelty
Suman Singh vs Sanjay Singh (8 Mar 2017)	Cruelty	Allowed	0.157	cruelty, divorce, mental
Narendra vs K Meena (6 Oct 2016)	Unspecified	Allowed	0.148	wife, divorce, husband
Jalendra Padhiary vs Pragati Chhotray (17 Apr 2018)	Cruelty, Desertion	Allowed	0.148	wife, cruelty, husband
Amutha vs A R Subramanian (19 Dec 2024)	Unspecified	Dismissed	0.108	husband, divorce

This is a real run of the actual TF-IDF + cosine-similarity pipeline in app.py against tagged_divorce_cases.csv (see caveat below).

3. Spot-check evaluation — precision@5

Four representative queries, each aimed at a specific ground of divorce, checked against whether the top-5 retrieved cases' ground_of_divorce field matched the intended theme:

Query theme	Precision@5	Top-5 verdicts
Cruelty	3/5 = 0.60	Allowed, Allowed, Allowed, Allowed, Dismissed
Mutual Consent	2/5 = 0.40	Allowed, Allowed, Allowed, Dismissed, Disposed
Desertion	3/5 = 0.60	Allowed, Allowed, Disposed, Allowed, Allowed
Irretrievable Breakdown	2/5 = 0.40	Allowed, Allowed, Allowed, Allowed, Allowed

Overall precision@5 across the 4 queries: 10/20 = 0.50

Reasonable, honest framing for the report: half of the top-5 results for a themed query carry a ground_of_divorce tag matching that theme. Given that 64% of cases are tagged "Unspecified," this actually undercounts true relevance — several "Unspecified" hits were topically on-point (shared vocabulary like "cruelty," "husband," "wife") but just weren't ground-tagged. Worth noting this nuance rather than the bare number alone.

So the corpus-side numbers (§1) are exact. The query-side numbers (§2, §3) are a very close approximation of what your deployed app would return — possibly off by a keyword or two in edge cases, but the ranking and verdicts should match closely.